

Jeeteshwar Shrivastava

jeet.shri.0909@gmail.com — linkedin.com/in/jeeteshwarshrivastava/ — +91 7563067047

Pune, India

Summary

Generative AI Engineer specializing in Large Language Models (LLMs), Natural Language Processing (NLP), and Transformer Architecture (BERT). Built scalable AI systems for intelligent automation using Prompt Engineering and Fine Tuning, deploying models handling 15,000+ daily inference requests. Experienced in full ML lifecycle (MLOps), system design, secure cloud-native AI platforms, and high-throughput ETL pipelines processing 2+ million records daily.

Experience

MotionCut, Remote

Jan 2024 – Feb 2024

Data Engineering Intern – AI Infrastructure & MLOps

- Implemented OAuth2 and JWT authentication for data pipelines and model deployment endpoints, reducing unauthorized access attempts by 28%.
- Accelerated data processing by 26% during model training using Redis and Memcached caching strategies.
- Parallelized AI inference with CUDA, increasing throughput by 30% and reducing inference latency by 22%.
- Reduced MLOps deployment failures by 12% through standardized testing, benchmarking, and model transparency checks.
- Optimized CI/CD pipeline with AWS, Kubernetes orchestration, and Git, cutting infrastructure costs by 18%.
- Re-architected data storage on AWS S3 using Apache Parquet with Spark and Presto, handling 120 GB of data.
- Automated CI workflows with Jenkins, reducing manual intervention by 6 hours weekly through automated testing and experiment tracking.

Projects

E-commerce Predictive Engine

- Built real-time ML inference architecture with FastAPI and SQLAlchemy, reducing data processing latency by 120 ms for rapid context modeling.
- Architected event-driven feature streaming pipeline with Apache Kafka, processing 1,000+ transactions per minute for analytics and NER.
- Secured feature store via OAuth2 role-based access control, blocking 100+ unauthorized manipulation attempts.
- Optimized PostgreSQL queries and ORM tuning for large-scale model optimization, boosting execution by 35%.
- Developed secure REST API gateways for MLOps integration, accelerating deployment velocity by 15%.

Tech Stack: Python, Scikit-Learn, FastAPI, SQLAlchemy, PostgreSQL, Apache Kafka, Celery, REST

Intelligent Content Platform

- Deployed deep learning application across GCP, AWS, and Vertex AI, optimizing GPU/CPU allocation for text generation, reducing infrastructure cost by 22%.
- Reduced HTTP request latency by 7% for PyTorch, TensorFlow, and Keras predictions via distributed Gunicorn and Uvicorn servers.
- Mitigated 12+ infrastructure vulnerabilities in data ingestion layer through comprehensive PyTest pipelines.
- Expedited unstructured data access via GraphQL APIs, MongoDB, and Redis caching, improving document retrieval by 32%.
- Automated data validation and continuous feedback workflows, accelerating foundation model training cycles by 25%.

Tech Stack: Python, Django, PyTorch, MongoDB, GraphQL, AWS, GCP, Docker, PyTest, MLflow

Technical Skills

- **AI, ML & NLP:** LLMs, Foundation Models, Transformers (BERT), RAG, Prompt Engineering, Semantic/Vector Search, Text Embeddings, Multi-agent Systems, Chatbots, Text Classification, Question Answering, Deep Learning, Transfer Learning, Neural Networks, Fine Tuning, Diffusion, Model Distillation, Compression.
- **Data Engineering & MLOps:** Big Data (Spark, SQL), ETL Pipelines, Distributed Systems, Databases, Feature Engineering, Model Optimization, Benchmarking, Experiment Tracking, Orchestration, CUDA.
- **Cloud & Frameworks:** GCP, AWS, Vertex AI, Kubernetes, Docker, Linux, Python, FastAPI, Flask, Django, GraphQL, REST APIs, TensorFlow, Keras, PyTorch, Git, CI/CD, Jenkins.

Certifications

- Deep Learning Specialization — deeplearning.ai / Coursera
- AWS Cloud Practitioner — Amazon Web Services (In Progress)

Core Competencies

- Break down complex algorithmic challenges into scalable ML architectures with strong analytical troubleshooting.
- Debug model-serving scripts and optimize enterprise-grade data pipelines with precision.
- Facilitate cross-functional execution on AI product development by bridging technical and non-technical stakeholders.
- Maintain clear engineering records across AI teams with rapid technology adoption.

Achievements

- Received “Mastermind Award” for engineering excellence and exceptional project documentation at MotionCut.

Education

Government Engineering College, Kaimur, Bihar
Bachelor of Technology in Computer Science Engineering
CGPA: 8.64

May 2026